

Healthcare With Artificial General Intelligence

Emmanuel Igwilo

ITAI 2372

Professor Devarakonda

April 22 2026

Artificial General Intelligence, abbreviated as AGI, represents one of the most transformative and debated concepts in the history of technology. AGI refers to a form of machine intelligence capable of performing any intellectual task that a human being can do. It is not constrained to a single domain, does not require retraining for new problems, and possesses the capacity to reason, learn, and adapt across an unlimited range of contexts.

Of all the industries that stand to be reshaped by AGI, healthcare offers perhaps the most compelling case. Medicine is pattern recognition, reasoning, and knowledge synthesis conducted under conditions of uncertainty, urgency, and human stakes. These are precisely the domains where a sufficiently capable general intelligence could operate with extraordinary effect. This case study examines the current state of healthcare, the role that narrow AI already plays within it, and the transformative potential that AGI could unlock, alongside the significant ethical and operational risks that must be addressed before such a transformation can responsibly occur.

Current State of Healthcare

Global healthcare is in a dire situation. In the United States, administrative costs consume an estimated 34 cents of every dollar spent on care. Physician burnout has reached epidemic proportions, with surveys consistently showing more than half of practicing doctors reporting symptoms of emotional exhaustion, depersonalization, or reduced sense of personal accomplishment. Diagnostic error rates remain stubbornly high, with studies estimating that roughly 12 million Americans are misdiagnosed in outpatient settings each year, and up to 40,000 patients in intensive care units may die annually from diagnostic mistakes.

Beyond diagnosis, the fragmentation of medical records across incompatible systems means that physicians routinely make decisions without access to a patient's complete history. Drug development timelines stretch to 10 to 15 years at costs exceeding one billion dollars per approved therapy, and the majority of clinical trials fail not because the science is wrong but because patient recruitment, protocol adherence, and data analysis processes are grossly inefficient. Meanwhile, global health inequity persists at scale: billions of people in low and middle income countries lack access to the specialist expertise that patients in wealthy urban centers take for granted.

Existing Uses of Narrow AI in Healthcare

Narrow AI has already made meaningful inroads across several areas of medicine. Computer vision models trained on medical imaging datasets now detect certain cancers, diabetic retinopathy, and cardiovascular anomalies with accuracy that rivals or exceeds that of experienced radiologists in controlled studies. Natural language processing tools extract structured data from clinical notes, reducing documentation burdens and enabling large scale retrospective research. Predictive algorithms flag patients at elevated risk for sepsis, readmission, or deterioration before clinical signs become apparent.

However, each of these systems is narrow by design and by limitation. A model that excels at reading chest radiographs cannot interpret the same patient's bloodwork, synthesize their social history, or reason about how a new symptom alters the differential diagnosis formed three weeks ago. Narrow AI assists; it does not integrate. It performs; it does not understand. The gap between what current AI can do and what clinical care actually demands is precisely the space that AGI could fill.

AGI Application

AGI differs from narrow AI in several fundamental respects. Where narrow AI is trained to optimize performance on a specific, predefined task using labeled data, AGI is capable of transferring learning across domains, reasoning under uncertainty about novel problems, generating hypotheses, and revising its own understanding in response to new evidence. It can engage in the kind of meta-cognitive reasoning that allows a skilled clinician to know not just what the answer is, but when they do not yet have enough information to answer, and what information would change their conclusion.

AGI is not simply a more powerful version of a large language model or a deeper neural network. It represents a qualitative shift in the nature of machine cognition: from pattern matching within a training distribution to genuine generalization beyond it. Whether and when such a system will exist is a matter of active scientific and philosophical debate. But the trajectory of current research, including advances in

reasoning architectures, multimodal integration, and autonomous learning, suggests that planning now for the healthcare implications of AGI is both prudent and necessary.

Use Cases in Healthcare

The most immediate and transformative application of AGI in healthcare would be the creation of persistent, longitudinal clinical intelligence systems. Unlike today's fragmented electronic health record platforms, an AGI-powered system would maintain a continuously updated, richly integrated model of each patient's health across their lifetime. It would synthesize genomic data, imaging, lab results, clinical notes, medication history, environmental exposures, and patient-reported outcomes into a coherent probabilistic picture that updates in real time. When a patient presents with a new complaint, the AGI would not search for the most statistically likely diagnosis in isolation. It would reason about how the new information fits with everything it already knows about that individual, flag inconsistencies, generate and rank differential diagnoses with explicit uncertainty estimates, and recommend targeted investigations.

In pharmaceutical research, AGI could compress drug development timelines dramatically. Today, the process of identifying a promising target, designing candidate molecules, predicting toxicity, designing trials, and analyzing results involves hundreds of specialists working across years. An AGI system could traverse this pipeline with far greater speed, generating novel hypotheses about disease mechanisms that no human researcher had considered, designing trials optimized for statistical power and patient burden simultaneously, and continuously reanalyzing accumulating trial data to detect efficacy signals or safety concerns earlier than traditional methods allow.

AGI could also begin to close the global specialist gap. A patient in rural areas in Africa with a complex neurological presentation currently has virtually no access to expert neurological assessment. An AGI diagnostic system accessible via a smartphone and a basic clinical examination, conducted by a community health worker, could provide the equivalent of specialist-level reasoning at the point of care. This is not merely a convenience. It represents a fundamental redistribution of medical expertise

that narrow AI, which requires curated infrastructure and domain-specific deployment, cannot achieve.

Benefits

The efficiency gains from AGI integration in healthcare are substantial and multidimensional. Administrative functions that currently consume a disproportionate share of clinical resources, including documentation, coding, prior authorization, and scheduling, could be handled autonomously, returning physician time to direct patient care. Diagnostic accuracy would improve, reducing the incalculable human and economic costs of missed or delayed diagnoses. Preventive interventions could be targeted with far greater precision, identifying patients at elevated risk before symptoms emerge and enabling earlier, cheaper, and more effective treatment.

Risks and Ethical Concerns

The deployment of AGI in healthcare carries risks that are not merely technical but deeply ethical and social. The most immediate concern is data privacy. An AGI clinical system by definition requires access to the most sensitive categories of personal data: genetic sequences, mental health records, reproductive history, financial information as it relates to health insurance, and behavioral data. The concentration of such information in systems controlled by governments or private corporations creates extraordinary risks of surveillance, discrimination, and exploitation. Robust legal frameworks, cryptographic protections, and patient ownership of data will be prerequisites, not afterthoughts.

Bias is a second major risk. Narrow AI systems have already demonstrated the capacity to perpetuate and amplify existing health disparities when trained on datasets that underrepresent certain populations. An AGI system with greater autonomy and broader reach could cause harm at greater scale if its training reflects historical inequities in who received care, what treatments were offered to whom, and whose outcomes were recorded. Deliberate, ongoing auditing for bias across demographic groups will be essential.

The workforce implications are also significant. Healthcare employs tens of millions of people worldwide, including many in roles that AGI could perform more efficiently. The transition will require active social policy, retraining programs, and a serious public conversation about how the productivity gains from AGI should be distributed. The goal must be a healthcare system that uses AGI to extend human capacity and reach, not to simply eliminate human labor while concentrating the resulting savings among institutional shareholders.

Conclusion

AGI holds the potential to be the most consequential technology in the history of medicine. By transcending the domain-specific limitations of narrow AI, it could address some of the deepest structural problems in healthcare: fragmented information, diagnostic error, inequitable access, and the paralyzing inefficiency of drug development. The vision of a system that reasons holistically about a patient's health across a lifetime, that democratizes specialist expertise, and that accelerates the discovery of new therapies is not science fiction. It is a plausible trajectory from where research currently stands.

Realizing that potential responsibly requires engaging seriously with the risks. Privacy protections, bias auditing, transparent governance, workforce transition support, and mechanisms for patient agency over their own data are not obstacles to AGI in healthcare. They are the conditions under which AGI in healthcare becomes something worth having. The most important decisions about this technology will not be made by engineers. They will be made by policymakers, clinicians, patients, and the public. Preparing those communities to participate meaningfully in those decisions is perhaps the most urgent work that can be done today.

Citations

Administrative Costs (34%) Himmelstein, D.U., Campbell, T., & Woolhandler, S. (2020). Health Care Administrative Costs in the United States and Canada, 2017. *Annals of Internal Medicine*, 172(2), 134–142. <https://pubmed.ncbi.nlm.nih.gov/31905376/>

Physician Burnout (~45%) Shanafelt, T.D. (2024). Changes in Burnout and Satisfaction With Work–Life Integration in Physicians and the General US Working Population Between 2011 and 2023. *Mayo Clinic Proceedings*.
[https://www.mayoclinicproceedings.org/article/S0025-6196\(24\)00668-2/fulltext](https://www.mayoclinicproceedings.org/article/S0025-6196(24)00668-2/fulltext)

12 Million Outpatient Misdiagnoses Annually Singh, H., et al. (2014). The Frequency of Diagnostic Errors in Outpatient Care: Estimations from Three Large Observational Studies Involving US Adult Populations. *BMJ Quality & Safety*, 23(9), 727–731.
<https://pmc.ncbi.nlm.nih.gov/articles/PMC4145460/>

40,500 ICU Deaths from Diagnostic Error Annually Winters, B., et al. (2012). Diagnostic Errors in the Intensive Care Unit: A Systematic Review of Autopsy Studies. *BMJ Quality & Safety*, 21(11), 894–902. <https://pubmed.ncbi.nlm.nih.gov/22822241/>

Drug Development Cost (~\$1–2.3 Billion) and Timeline (10–15 Years) Wouters, O.J., McKee, M., & Luyten, J. (2020). Estimated Research and Development Investment Needed to Bring a New Medicine to Market, 2009–2018. *JAMA*, 323(9), 844–853.
<https://pubmed.ncbi.nlm.nih.gov/32125404/>

Deloitte (2022). Biopharmaceutical R&D Cost Report. Summarized at:
<https://www.genengnews.com/gen-edge/the-unbearable-cost-of-drug-development-deloitte-report-shows-15-jump-in-rd-to-2-3-billion/>

Congressional Budget Office (2021). Research and Development in the Pharmaceutical Industry. <https://www.cbo.gov/publication/57126>