

ARTIFICIAL INTELLIGENCE IN GOVERNMENT

By Emmanuel Igwilo

ITAI 2372

Artificial intelligence is fundamentally reshaping how governments deliver services, protect public resources, and make decisions that affect millions of citizens. From machine learning algorithms that detect fraudulent benefits claims to computer vision systems that monitor infrastructure health, AI technologies are increasingly embedded in the operational fabric of public institutions worldwide.

The scale of this transformation is significant. According to the OECD, more than 60 countries have now adopted national AI strategies, and government agencies at every level are piloting or deploying AI tools across domains as varied as tax administration, emergency management, social services, and urban planning. The global public sector AI market was valued at approximately \$8.4 billion in 2024 and is projected to exceed \$24 billion by 2030, reflecting an accelerating commitment to data-driven governance.

Problems

Medicare and Medicaid fraud represents one of the most costly forms of government waste in the United States, costing taxpayers an estimated \$60 to \$90 billion annually. Prior to the introduction of AI-based systems, the Centers for Medicare and Medicaid Services (CMS) relied on a predominantly pay-and-chase model: claims were processed and paid, and investigators subsequently attempted to recover funds from fraudulent providers. This reactive approach was chronically underfunded relative to the scale of abuse and resulted in billions of dollars in irrecoverable losses each year.

The problem was structural. With over one billion claims processed annually, human reviewers could only meaningfully scrutinize a tiny fraction of transactions. Fraudulent schemes were becoming increasingly sophisticated, exploiting regulatory gaps and mimicking the billing patterns of legitimate providers. A scalable, proactive solution was urgently required.

AI Tools

In 2011, CMS deployed the Fraud Prevention System (FPS), developed in partnership with Northrop Grumman and later enhanced by IBM Watson analytics capabilities. The system employs several interconnected AI and machine learning techniques:

AI Technique	Application in Fraud Detection
Supervised Learning	Trained on historical data of confirmed fraud cases to classify new claims as high, medium, or low risk.
Unsupervised Anomaly Detection	Identifies billing patterns that deviate statistically from peer providers without requiring pre-labeled examples.
Social Network Analysis	Maps relationships between providers, suppliers, and patients to identify coordinated fraud rings.
Predictive Scoring	Assigns real-time risk scores to claims, enabling pre-payment suspension of suspicious transactions.
Natural Language Processing	Parses clinical documentation to flag inconsistencies between diagnoses, procedures, and billing codes.

The system analyzes over 4.5 million claims per day in near real time, generating risk scores that allow investigators to prioritize their limited review capacity on the highest-probability fraud cases.

Benefits

The results of the FPS deployment have been substantial and well-documented. The following summary table captures key performance metrics reported by CMS and the Department of Health and Human Services Office of Inspector General:

Performance Metric	Reported Outcome
Total fraud savings (FY2011 to FY2023)	Approximately \$42 billion identified or recovered
Return on investment	\$10 to \$12 saved per \$1 invested in the system
Increase in fraud referrals	Over 2,500% increase in actionable leads for investigators
Pre-payment stops (FY2022)	>\$1.3 billion in fraudulent claims halted before disbursement
Provider exclusions supported	More than 1,100 high-risk providers removed from the program

Beyond financial recovery, the system has shifted CMS from a reactive posture to a proactive one, allowing investigators to interrupt fraud schemes before payments are made rather than pursuing costly legal recovery after the fact. It has also enabled more targeted audits, reducing burden on legitimate providers who would otherwise face indiscriminate review.

Challenges

Despite its success, the CMS FPS also illustrates the persistent challenges of AI deployment in government contexts:

- **Algorithmic Bias:** Early versions of the system produced disproportionate false-positive rates among certain provider categories, including rural physicians and minority-owned practices, raising concerns about discriminatory enforcement patterns.
- **Explainability:** The black-box nature of some machine learning models made it difficult for investigators to explain why specific providers were flagged, creating legal vulnerabilities in enforcement proceedings that require evidentiary justification.
- **Data Quality:** The system depends on the integrity of claims data, which varies significantly across states and managed care organizations, requiring substantial ongoing data governance investment.

Innovations

First, infrastructure failure is both predictable and preventable. Research from the American Society of Civil Engineers and MIT's Urban Planning Lab consistently demonstrates that up to 70% of infrastructure failures show detectable warning signs in sensor and inspection data months before catastrophic failure occurs. Current maintenance regimes typically fail to act on this data due to capacity constraints and fragmented information systems.

Second, infrastructure degradation is inequitably distributed. Studies of cities including Detroit, Baltimore, and Chicago have documented systematic patterns in which lower-income and predominantly minority neighborhoods receive less proactive maintenance, experience more service disruptions, and wait longer for repairs when failures occur. AI can function as a depoliticizing force by grounding allocation decisions in objective risk metrics rather than political influence or administrative discretion.

Third, the technology is ready. Sensor costs have dropped by more than 80% over the past decade, making city-wide IoT deployments economically feasible. Proven machine learning architectures for predictive maintenance exist in the industrial sector, where similar approaches have reduced unplanned downtime in manufacturing and energy infrastructure by 30 to 50%.

System Architecture and AI Methods

Component	Technology	Function
Sensor Network	IoT sensors, accelerometers, acoustic monitors	Continuous structural health monitoring
Data Integration Layer	Cloud data lake, real-time stream processing	Consolidate sensor, GIS, weather, complaint data
Predictive Engine	LSTM neural networks, gradient boosting	Failure probability scoring by asset and location
Equity Module	Fairness-constrained optimization algorithm	Adjust maintenance scheduling for equity goals
Citizen Interface	Mobile app, open API, 311 integration	Transparent reporting and community input
Audit Dashboard	Explainable AI (XAI) visualization tools	Accountability and bias monitoring for officials

A critical design principle of this proposal is the equity module, which applies fairness-constrained optimization to maintenance scheduling. Rather than simply maximizing aggregate system uptime, the algorithm is designed to minimize disparities in service quality across neighborhoods, treating equity not as a secondary consideration but as a primary optimization target embedded in the model objective function.

Based on analogous deployments in Kansas City, Singapore, and Amsterdam, projected outcomes within a five-year implementation horizon include:

- A 35 to 45% reduction in unplanned infrastructure failures citywide, with disproportionate improvements in previously underserved areas.
- An estimated 20 to 30% reduction in total infrastructure maintenance costs through proactive intervention, which is typically three to seven times cheaper than emergency repair.
- Measurable narrowing of service quality gaps between high-income and low-income neighborhoods, as tracked through an open public equity dashboard.
- Improved citizen trust and engagement through transparent, explainable AI-driven decision-making made accessible via a public-facing app.
- A replicable governance model that other municipalities can adopt, with potential to be scaled nationally through federal infrastructure partnerships.

Potential Challenges:

- **Data Infrastructure Gaps:** Many municipalities, particularly smaller cities and those in lower-income regions, lack the baseline digital infrastructure to support large-scale IoT sensor deployment. Federal co-investment and phased implementation plans will be essential.
- **Algorithmic Accountability:** Without robust governance frameworks, AI-driven allocation could become a new mechanism for inequity if equity metrics are poorly defined or if the system is manipulated. Independent auditing and community oversight boards are non-negotiable design requirements.
- **Workforce Transition:** The shift from reactive to predictive maintenance requires significant retraining of municipal workforce and may face resistance from labor organizations concerned about job displacement.
- **Vendor Lock-in Risk:** Reliance on proprietary AI platforms creates long-term dependency risks. Procurement policies should favor open-source or interoperable architectures wherever feasible.
- **Privacy and Surveillance Concerns:** Sensor networks that monitor urban activity create legitimate surveillance concerns that must be addressed through strict data minimization policies and independent oversight.

Conclusion

The CMS Fraud Prevention System case study demonstrates both the substantial potential and the genuine complexity of AI deployment in the public sector. When implemented thoughtfully, AI can generate extraordinary returns on public investment, protect taxpayer resources, and shift government agencies from reactive to proactive operating modes. The \$42 billion in identified fraud savings and the 2,500% increase in actionable investigative leads represent transformative outcomes that would be inconceivable through conventional approaches.

The proposed smart city infrastructure management system represents a logical and timely extension of proven AI capabilities into one of the most tangible and consequential domains of local governance. Infrastructure is the physical substrate of civic life, and its equitable maintenance is a prerequisite for any meaningful claim to public fairness. By embedding equity as a primary objective function rather than an afterthought, this proposal advances a vision of AI in government that serves not just efficiency, but justice.

References

Department of Health and Human Services, Office of Inspector General. (2023). Annual report to Congress on the Medicare and Medicaid Anti-Fraud programs. U.S. Government Publishing Office.

American Society of Civil Engineers. (2025). 2025 Report Card for America's Infrastructure. ASCE. <https://www.infrastructurereportcard.org>

OECD. (2024). OECD AI Policy Observatory: National AI strategies and policies. Organisation for Economic Co-operation and Development. <https://oecd.ai/en/dashboards/overview>

Obermeyer, Z., Powers, B., Vogeli, C., and Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447 to 453. <https://doi.org/10.1126/science.aax2342>

Raji, I. D., and Buolamwini, J. (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society. <https://doi.org/10.1145/3306618.3314244>

Smart Cities World. (2024). Predictive infrastructure maintenance: Case studies from Singapore and Amsterdam. Smart Cities World Annual Review. <https://www.smartcitiesworld.net>

U.S. Government Accountability Office. (2022). Artificial intelligence in government: Agencies have taken steps to use AI, but governance mechanisms need strengthening (GAO-22-104880). GAO. <https://www.gao.gov/products/gao-22-104880>

World Economic Forum. (2024). AI governance in the public sector: Principles, frameworks, and emerging practice. WEF White Paper. <https://www.weforum.org/reports/ai-governance-public-sector>

Zheng, S., Trott, A., Srinivasa, S., Naik, N., Gruesbeck, M., Parkes, D. C., and Socher, R. (2022). The AI Economist: Improving equality and productivity with AI-driven tax policies. arXiv:2004.13332. <https://arxiv.org/abs/2004.13332>