

Artificial Intelligence Ethics and Societal Implications

Case Study: Facial Recognition in Law Enforcement

Emmanuel Igwilo

ITAI 2372

April 4, 2026

Introduction

Artificial intelligence is reshaping critical public systems, and few applications raise more urgent ethical questions than facial recognition technology (FRT) used in law enforcement. Deployed across police departments, border agencies, and surveillance networks worldwide, FRT promises to enhance public safety and streamline criminal investigations. Yet it also introduces serious concerns around civil liberties, racial bias, and privacy. This report examines FRT in law enforcement as a case study in the complex relationship between technological innovation and ethical responsibility, analyzing its benefits, ethical challenges, societal effects, risks, and potential remedies.

Overview of the AI System

Facial recognition technology in law enforcement refers to the automated identification of individuals by analyzing facial features extracted from images or video footage. Agencies use FRT to match surveillance camera captures against criminal databases, identify suspects from crime scene photos, verify identities at border checkpoints, and monitor public spaces in real time.

At a high level, FRT works through a multi-stage pipeline. A camera captures a face, a detection algorithm locates and crops it, and a deep neural network converts the face into a numerical vector called a faceprint. A matching algorithm then compares this vector against a reference database; if the similarity score meets a set threshold, the system returns a candidate match. The core problem FRT aims to solve is the slow, manual process of identifying unknown individuals from images. Traditional methods relied on human analysts reviewing mugshots, which could take days. FRT can return results in seconds. Real-world deployments include the FBI's Next Generation Identification system, the New York Police Department's use of Clearview AI, and the Metropolitan Police Service in London's live FRT trials in public areas.

Benefits and Positive Impact

The most significant operational advantage of FRT is speed. Investigations that once took analysts days can now yield results in seconds, which can be critical in time-sensitive cases. Under favorable conditions, such as high-resolution frontal images and good lighting, modern FRT systems can achieve identification accuracy rates exceeding 99%, as benchmarked by the National Institute of Standards and Technology (NIST). This level of accuracy provides investigators with leads that human analysts might take far longer to develop.

Automating the initial phase of suspect identification also reduces labor costs, helping under-resourced agencies process large volumes of images without proportional increases in staffing. Beyond criminal investigation, FRT has been used to identify missing children, locate trafficking victims, and identify deceased individuals in disaster scenarios. Interpol operates a facial recognition database used in cross-border trafficking and child exploitation investigations, where FRT has contributed to documented rescues.

Ethical Issues

Racial Bias and Algorithmic Discrimination

Independent research has consistently shown that commercial FRT systems perform with significantly different accuracy rates across demographic groups. A 2019 audit by NIST found that many widely used algorithms had error rates up to 100 times higher for Black and Asian faces compared to white faces and performed worse on women than men. Because these algorithms are trained on datasets that over-represent lighter-skinned male faces, they encode existing demographic disparities into their outputs.

The real-world consequences are severe. In at least three documented cases in the United States involving Robert Williams, Michael Oliver, and Nijeer Parks, Black men were wrongfully arrested based on erroneous FRT matches. In each case, the algorithm produced a false positive that investigators accepted without adequate scrutiny. These are documented miscarriages of justice directly traceable to biased AI systems operating in high-stakes contexts.

Privacy and Mass Surveillance

Live FRT systems scan the faces of everyone who passes through a monitored area, the vast majority of whom are innocent bystanders with no connection to any investigation. This constitutes a form of blanket surveillance that challenges foundational norms around privacy, freedom of movement, and the presumption of innocence. Critics including the ACLU and the European Data Protection Board argue that mass FRT deployment is incompatible with constitutional and human rights protections. The aggregation of location data associated with FRT logs also enables the reconstruction of individuals' movements and associations without their consent or a warrant.

Lack of Transparency and Accountability

Most commercial FRT systems are proprietary, meaning law enforcement agencies cannot independently audit the algorithms they rely on, and defense attorneys cannot fully challenge the basis of a match in court. This opacity undermines due

process rights and makes meaningful oversight very difficult. When a false positive result contributes to an arrest, there is often no mechanism for the affected individual to understand or contest the algorithmic basis of the accusation, representing a fundamental accountability gap that has not been adequately addressed.

Societal Implications

Because FRT systems perform least accurately on Black, Asian, and female faces, and because these communities are already disproportionately subject to policing, the technology risks automating and amplifying existing racial disparities in the criminal justice system. Biased AI tools embedded in institutional processes can acquire a veneer of scientific objectivity that makes discrimination harder to challenge.

The awareness of pervasive FRT surveillance can also suppress lawful activity. Research in behavioral psychology suggests that surveillance awareness reduces willingness to participate in political protests or other monitored activities, creating a chilling effect on democratic participation. Public attitudes toward FRT in policing are divided: support drops sharply when respondents are informed about error rates, particularly for minority populations. Several major U.S. cities, including San Francisco, Boston, and Portland, have enacted bans on government use of FRT in response to public opposition.

Around the world, FRT deployment varies enormously by national context. China's Social Credit System represents the most extensive integration of FRT into state governance, enabling real-time monitoring of hundreds of millions of people. Authoritarian states can also export FRT infrastructure, as Chinese firms have done across parts of Africa and Central Asia, potentially enabling surveillance regimes that suppress political dissent far beyond their original development contexts.

Risks and Challenges

Despite strong benchmark performance, FRT accuracy degrades in real-world conditions: poor lighting, obscured faces, aging, low-resolution footage, and changes in pose can all substantially increase error rates. The gap between controlled performance and operational performance is frequently underestimated by agencies deploying FRT.

Biometric databases holding millions of faceprints are high-value targets for malicious actors. Unlike passwords, facial biometrics cannot be changed if compromised. A large-scale breach, similar to the 2019 U.S. Customs and Border Protection incident, could expose individuals to identity fraud and targeted surveillance with no practical remedy. Law enforcement tools also have a history of expanding beyond their originally justified use cases. FRT deployed for counterterrorism investigations may later be applied to minor offenses or political protest monitoring. Without firm legal constraints on use cases, scope creep is difficult to prevent. Finally, the patchwork of local bans, pilot programs, and voluntary guidelines currently governing FRT creates legal uncertainty and no clear minimum standards for accuracy, bias testing, or oversight.

Recommendations

Addressing the risks of facial recognition in law enforcement requires coordinated action across policy, technical, and institutional dimensions.

Policy and Regulation

- Binding accuracy and bias standards: Enact legislation establishing minimum accuracy thresholds, mandatory demographic bias auditing, and independent certification before deployment.
- Transparency requirements: Require agencies to publish use policies and audit logs annually and disclose FRT evidence to defendants in criminal proceedings.
- Restrict real-time surveillance: Prohibit FRT use in live public monitoring without a judicial warrant based on individualized suspicion.
- Use-case limitations: Establish strict legal boundaries to prevent FRT deployed for serious crime investigation from being repurposed for minor offenses or immigration enforcement.

Technical Improvements

- Demographic parity testing: Require vendors to demonstrate equitable accuracy across race, gender, and age groups before deployment, with periodic re-auditing.
- Human-in-the-loop requirements: Mandate that FRT outputs be treated as investigative leads requiring independent corroboration, never as standalone grounds for arrest.

Oversight and Accountability

- Independent oversight bodies: Establish multi-stakeholder boards with authority to audit deployments and issue binding recommendations.
- Community engagement: Require formal public consultation processes and published impact assessments before any agency deploys FRT.

Conclusion

Facial recognition technology in law enforcement illustrates the dual nature of artificial intelligence: a tool of genuine operational utility that simultaneously introduces serious risks to justice, equality, and civil liberty. The efficiency gains are real, but so are the documented harms, including wrongful arrests, racially discriminatory error rates, and accountability gaps that undermine due process. Responsible deployment requires binding accuracy standards, demographic equity requirements, strict use-case limitations, and robust independent oversight. Without these safeguards, FRT risks encoding and amplifying the very injustices that effective law enforcement should exist to prevent. The choices societies make now about regulating these systems will shape the scope of civil liberties for decades to come.

References

- Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research*, 81, 1-15.
- Grother, P., Ngan, M., & Hanaoka, K. (2019). *Face recognition vendor test (FRVT) part 3: Demographic effects*. National Institute of Standards and Technology.
<https://doi.org/10.6028/NIST.IR.8280>
- Hill, K. (2020, June 24). Wrongfully accused by an algorithm. *The New York Times*.
<https://www.nytimes.com/2020/06/24/technology/facial-recognition-arrest.html>
- Metropolitan Police Service. (2022). *Live facial recognition: Oversight board annual report*.
<https://www.met.police.uk/police-data/facial-recognition>
- Raji, I. D., & Buolamwini, J. (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. *Proceedings of AAAI/ACM Conference on AI, Ethics, and Society*, 429-435.
- American Civil Liberties Union. (2021). *The perpetual lineup: Unregulated police face recognition in America*. ACLU.
<https://www.aclu.org/issues/privacy-technology/surveillance-technologies/face-recognition-technology>